

# Evaluating and developing Monero security in a post-quantum world

Adam Corbo, Mitchell Krawiec-Thayer (@isthmus)\*, Brandon G Goodell†

Insight, Monero Research Lab

Aug 2020

## Abstract

This technical note renders a thorough analysis of Monero’s extant weaknesses before a theoretical quantum adversary. Extrapolating from its survey of weaknesses, the discussion herein extends into technical descriptions of plausible solutions, with a focus on their respective practical and theoretical trade-offs. [Revisit this when nearing completion](#)

## 1 Introduction

### 1.1 Prerequisites and preliminaries

We denote concatenation of strings with  $||$ . For a finite unordered set  $X$  with  $|X| = n$ , note that we can arbitrarily label entries in  $X$  and presume  $X = [n] = \{1, 2, \dots, n\}$  without losing any generality.

A classical bit is a binary digit that takes on one of two values to indicate a logical value. By convention, these are 0 and 1. Qubits, on the other hand, represent a spectrum of logical values that reduces to a classical bit whenever the qubit interacts with the environment. Indeed, to measure the state of a qubit is equivalent to collapsing it into a classical state. Before measurement, a qubit remains in a superposition of the values 0 and 1. States of qubits are elements of  $\mathbb{C}^2$ ; we denote orthonormal basis vectors  $|1\rangle$  and  $|0\rangle$ . The superposition  $|\psi\rangle$  of a qubit is represented as a linear combination of these basis vectors  $|\psi\rangle = a_0|0\rangle + a_1|1\rangle$  where  $a_0 \in \mathbb{C}$  is the complex scalar amplitude of the state along the direction of  $|0\rangle$  in  $\mathbb{C}^2$ , and  $a_1$  is the amplitude along  $|1\rangle$  in  $\mathbb{C}^2$ .

Amplitudes may be thought of as “quantum probabilities:” the amplitude along an orthonormal basis vector is related to the probability that the collapsed state corresponds to that vector. That is to say, the amplitude along  $|0\rangle$  is related to the probability that the collapsed state is 0, and the amplitude along  $|1\rangle$  is related to the probability that the collapsed state is 1. We must take care, however. Amplitudes are represented by complex numbers, while traditional probabilities are described by real numbers. Just as the probabilities of a classical system must integrate to 1 under some probability measure  $\mu$  in order to form a distribution function, the squared magnitudes of state amplitudes in a quantum system must satisfy  $\int |a_0|^2 + |a_1|^2 d\mu = 1$ .

Whereas a classical adversary, which we abbreviate with CA, uses computers that operate only with bits, a theoretical quantum adversary (abbreviated QA) has a computer that operates with qubits. The CA adversary can only be expected to complete algorithms using probabilistic Turing machines in polynomial time (PPT algorithms) whereas the QA adversary can solve bounded error, quantum, polynomial (BQP) time algorithms. The algorithms from Sections 1.2.1, 1.3, and 1.4 are BQP algorithms and require qubits to execute efficiently, these algorithms when run with qubits solve certain problems exponentially faster than any CA.

We say two transactions are *unlinkable* if it is difficult for a non-recipient to discern whether the two transactions have the same recipient or not. We say transactions are *signer-ambiguous* if it is difficult for a non-sender to discern the sender of the transaction. We say transactions are *confidential* if it is difficult for a non-sender, non-recipient to discern the transaction amounts.

---

\*Isthmus@getmonero.org

†surae@getmonero.org

## 40 1.2 The Quantum Adversary’s Capabilities

### 41 1.2.1 Violate Discrete Logarithm Hardness with Shor’s Algorithm

42 For a finite abelian group  $G$ , a subgroup  $H \subseteq G$ , a finite set  $X$ , and a function  $f : G \rightarrow X$ , we say  $f$  *hides*  $H$  if,  
43 for any  $g_1, g_2 \in G$ ,  $f(g_1) = f(g_2)$  if and only if  $g_1H = g_2H \in G/H$ . For example, setting  $X$  as the quotient group  
44  $X = G/H$ , we see that the canonical group epimorphism  $f : G \rightarrow G/H$  hides  $H$ . We can always answer oracle  
45 queries made to  $f$  with  $O(\log |G| + \log |X|)$  bits. Before Shor’s algorithm, it was thought that the following hardness  
46 assumption was valid.

47 **Assumption 1.1** (Hidden Subgroup Problem). *There does not exist an algorithm  $\mathcal{A}$  with oracle access to  $f$  that*  
48 *can output a generating set for  $H$  in polynomial time.*

49 Shor’s algorithm for discrete logarithms is a BQP algorithm capable of solving the hidden subgroup problem,  
50 violating Assumption 1.1. It is still thought that Assumption 1.1 is valid when restricted to PPT algorithms, such  
51 that there does not exist a PPT algorithm that violates Assumption 1.1.

52 Indeed, if  $G' = \langle g \rangle$  is any group of order  $p$ , we can compute the discrete logarithm of an ostensibly random  $X = g^x$   
53 using  $G = \mathbb{Z}_{p-1}^2$ ,  $H = \langle (X, g^{-1}) \rangle$ , and the function  $f : G \rightarrow \mathbb{Z}_p^*$  defined by mapping  $(a, b) \mapsto g^a X^b$ . This  $f$  is a group  
54 homomorphism whose kernel is  $H$ ; finding  $H$  is equivalent to computing the discrete logarithm for the generator  
55  $X = g^x$ . A more in-depth technical outline for running Shor’s algorithm for breaking the discrete logarithm hardness  
56 assumption is outlined in Section A.1. For more details, see [10] and [6].

## 57 1.3 Utilize Grover’s Algorithm to Find Pre-Images of Hash Functions and Unstruc- 58 tured Search Capabilities

### 59 1.3.1 Grover’s Algorithm

60 Grover’s algorithm can be applied to compute hash digest pre-images: this algorithm finds unordered database entries  
61 that satisfy search criteria in  $O(\sqrt{n})$  time, where  $n$  is the database size. See [4] for details, since we do not specify the  
62 algorithm here. It is shown in [1] that Grover’s algorithm is asymptotically optimal even for quantum adversaries.

63 Let  $f : [n] \rightarrow \{0, 1\}$  be a function which describes whether an index matches the search criteria and to which  
64 Grover’s algorithm has oracle access. For any non-negative function  $f : [n] \rightarrow \mathbb{Z}^+$  such that  $\sum_{x \in [n]} f(x) > 0$ , note  
65 that non-negativity implies there must exist some  $w \in [n]$  such that  $f(w) > 0$ . Moreover, the codomain of  $f$  is  $\{0, 1\}$   
66 so  $f(w) = 1$ .

67 Grover’s algorithm makes approximately  $O(\sqrt{n})$  queries to  $f$  via oracle access and outputs a solution  $w$  such that  
68  $f(w) = 1$ . Given a hash function  $H$ , we can define  $f(x)$  to be 1 when  $H(x) = y$  and 0 otherwise and use Grover’s  
69 algorithm to find hash digest pre-images that fit the necessary parameters.

70 In many cases, the universe of possible pre-images is stored on the blockchain in public, and so applying Grover’s  
71 algorithm in these cases in  $O(\sqrt{N})$  time can be quite fast. Grover’s can still be applied to find arbitrary pre-images  
72 that have not yet necessarily been posted publicly, but the search space is significantly larger. In these cases,  $O(\sqrt{N})$   
73 is often still a prohibitive hurdle.

74 For a technical description of Grover’s algorithm, see Section A.2.

### 75 1.3.2 Difficulty of Finding the Pre-Image of Hash Digests

76 Some aspects of security rely on the fact that finding the pre-image of a hash digest is difficult. For a QA, as stated  
77 in the last subsection, employing Grover’s algorithm, it is possible to find a marked value in an un-ordered set of  
78 data in approximately the square root of the number of possible entries to look through. One possible mitigation for  
79 this is to double the key lengths used in the hash function, which makes this more difficult to accomplish.

80 To brute force a 256 bit collision resistant hash digest, essentially trying every possible input pre-image to a hash  
81 function until the known hash digest gets spit out, on a classical computer would take on average  $2^{255}$  operations.

82 This kind of attack, with the state space of about  $2^{256}$  is highly impractical to attempt and as such is considered one  
83 way.

84 Using Grover's algorithm, one could attempt this same attack by setting up the Oracle circuit used in Grover's  
85 algorithm with the same hash function in order to mark when the appropriate pre-image would spit out the known  
86 hash digest. If you take into account the state space of all possible inputs for a 256 bit collision resistant hash function,  
87 this would still take approximately  $2^{128}$  operations of Grover's algorithm to generate an output that corresponded  
88 with the appropriate pre-image for the known hash digest. While still significantly better at attempting this attack  
89 than a classical computer, for standard clock speeds this would still require a time period greater than the age of the  
90 universe to attempt.

91 If the state space is lowered, for instance if we know one or more of the factors used to generate the hash digest,  
92 or can narrow down the possible pre-images in some way, then this attack becomes more feasible for a quantum  
93 computer to attempt. Depending on how effective this is at reducing the search space, it will still in most cases take  
94 longer than it takes Shor's to break ECC, RSA, etc.

95 Grover's algorithm can also be used in conjunction with other methods, such as QDC (Quantum Differential  
96 Cryptanalysis), to find the pre-image of hash digests. (See 1.4.2).

## 97 1.4 Simon's Algorithm, Quantum Differential Cryptanalysis, and Further Capabili- 98 ties.

### 99 1.4.1 Simon's Algorithm

100 Simon's algorithm, from [13], can be applied to extract XOR masks from functions under which they are invariant.  
101 Since many hash functions are based on iterated XOR masks, this allows for differential cryptanalysis.

102 Let  $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$  be any function that is invariant under some mask  $a$ . That is to say, there exists some  
103  $a \in \{0, 1\}^n$  such that, for every  $x, y \in \{0, 1\}^n$ ,  $f(x) = f(y)$  if and only if  $x \oplus y \in \{0, a\}$ . Given oracle access to  
104  $f$ , Simon's algorithm makes  $O(n)$  queries to the oracle and produces as output the mask  $a$ . It is shown in [5] that  
105 Simon's algorithm is asymptotically optimal.

106 Simon's algorithm can be used to set up a system of linear equations that can be used to find the outputs and find  
107 the XOR mask of certain functions, which has numerous applications in cryptography. For a technical description  
108 of Simon's algorithm, see Section A.3.

### 109 1.4.2 Quantum Differential Cryptanalysis

110 Besides brute force attacks, there are some other methods a QA could leverage to reverse or reveal hidden information  
111 about the pre-image of a hash digest. Using Simon's algorithm, it's possible to create a system of linear equations  
112 that can be used to perform differential cryptanalysis to decrypt an XOR mask, or attack symmetric key primitives  
113 [12]. For block ciphers relying on a Feistel scheme, which if they are used in the construction of a hash function,  
114 would be vulnerable to this advantage a QA could leverage.

115 Attacks of this nature, namely using an efficient quantum algorithm (Such as the Bernstein-Vazirani algorithm  
116 [14], which is beyond the scope of this technical note) to set up a system of linear equations, and then solving these  
117 classically and use them to perform differential cryptanalysis, are a bit more difficult to describe than the brute force  
118 Grover's algorithm attack. Attacks of this nature require a deep case by case look into the internal architecture  
119 of each specific hash function used, (Keccak, chacha20 in the case of Monero) Besides Simon's algorithm and the  
120 Bernstein-Vazirani algorithm, other algorithms exist that could be used to perform differential cryptanalysis on the  
121 block ciphers used in the construction of each hash function.[?]

122 As far as this technical audit could surmise, Keccak (the hash function used in Monero) is secure against currently  
123 known methods employing Simon's or the Bernstein-Vazirani algorithm to perform differential cryptanalysis and  
124 find information about or reverse the pre-image of the Keccak hash function. That being said, there could be other

125 methods using these algorithms that were not covered during this audit, with or without using Grover’s algorithm  
126 in tandem.

127 For that matter, there are other algorithms as well that can be used to perform quantum differential cryptanalysis  
128 other than B-V or Simon’s, it is unknown if others exist that aren’t as well known as of this writing.

129 Finding or designing a hash function whose internal architecture is provably resistant to all possible methods  
130 of Quantum Differential Cryptanalysis (henceforth QDC), is outside the scope of this technical audit. It could be  
131 possible Keccak is already secure against all possible forms of QDC, but until such a proof is found, we will assume  
132 for the purposes of this audit that it might be possible for a future QA to utilize such an advantage in an extreme  
133 case, maybe in combination with Grover’s algorithm, to find the pre-image of a known hash digest.

134 To summarize, pure brute force attacks on Keccak using Grover’s is intractable for a QA. However, by decreasing  
135 the state space this attack can become possible to implement for a QA if not a CA, though possibly still difficult.  
136 There are a few methods to do this, which could include QDC. Since it is unknown whether or not keccak specifically  
137 is susceptible to QDC, it might still be possible for a QA to find the pre-image of a known hash digest even without  
138 decreasing the possible state space.

### 139 1.4.3 Further Capabilities

140 The capabilities of a QA as outlined in this article primarily center around the advantages that can be leveraged  
141 by independently utilizing Shor’s, Grover’s, Simon’s or some combination of these algorithms to attack the security  
142 features present in Monero. The further methods of QDC are not extensively explored, such as using Simon’s or the  
143 B-V algorithm, but also should not be considered a complete list of algorithms one could attempt QDC with. A QA  
144 also inherently has the added capability of being able to generate true randomness rather than relying on a PRNG,  
145 which potentially could be useful as well.

## 146 2 Technical Overview of Vulnerabilities

147 For the purposes of probing potential vulnerabilities, we assume a theoretical QA capable of efficiently leveraging  
148 the above-detailed quantum-empowered algorithms (as well as, prospectively, other algorithms not yet discovered).  
149 If such an adversary were to exist, a number of Monero’s core mechanisms would be vulnerable to the plausible  
150 implementation of such algorithms. Below, we describe how these mechanisms are impacted by various known  
151 algorithms. The following should not be considered a comprehensive list.

### 152 2.1 Deriving Wallet Seeds

153 In this section, we explain how an efficient QA can derive wallet seeds from public information like addresses, sub-  
154 addresses, and data intended to be posted publicly on the blockchain like pairs of one-time keys with the same  
155 recipient.

156 First, we recall a bit about the key generation process in Monero. For a group  $G$  of order  $p$  with generator  $g$ ,  
157 a Monero-style wallet generation involves two keypairs, one for spending and one for viewing. However, the private  
158 spend key  $k_s$  is used to deterministically compute all other keys in the Monero Core wallet.

159 The private spend key  $k_s \in \mathbb{Z}_{p-1}$  is a random integer sampled from a pseudorandom number generator. The  
160 public spend key  $K_s$  is a group element obtained by computing  $K_s = g^{k_s}$ . The 25-word mnemonic “seed phrase”  
161 used to backup a wallet is simply  $k_s$  (with a checksum) encoded into a base-1626 dictionary for convenience.

162 A private view key  $k_v$  is generated and the corresponding public key  $K_v = g^{k_v}$  is computed. In the Monero  
163 Core wallet, the private view key is derived from the private spend key by computing  $k_v = H(k_s)$ . Further details  
164 are unnecessary for our analysis, and it should be noted that  $k_s$  and  $k_v$  could be independently generated with a  
165 pseudorandom number generator to prevent a QA from exploiting this link.

166 The wallet’s primary public address  $A$  is the base-58 encoded concatenation of a network prefix  $N$ , both public  
167 keys  $K_v$  and  $K_s$ , and a checksum  $C$ . Specifically,  $A = N||K_s||K_v||C$ .

### 168 2.1.1 Key Extraction From Addresses

169 The public spend key  $K_s$  and the public view key  $K_v$  can be parsed directly from the address  $A$ , enabling a QA that  
170 learns of any Monero address to apply Shor’s algorithm to extract the corresponding private spend key  $k_s$ . From  
171 this, a CA can compute  $k_v$  (if  $k_v$  is deterministically derived from  $k_s$  as in the Core implementation). Even if a user  
172 is not using a view key  $k_v$  deterministically derived from  $k_s$ , a user who is using two independent pseudorandom  
173 numbers for  $k_s$  and  $k_v$  is still vulnerable to the QA, who can compute  $k_v$  by inverting the map  $k_v \mapsto g^{k_v} = K_v$  using  
174 Shor’s algorithm a second time.

175 The adversary then essentially owns the wallet: they can derive the remaining keys, view the entire history of  
176 the wallet, spend any funds within, and so on. In this way, even publishing a public key (i.e. posting your address)  
177 could be dangerous.

### 178 2.1.2 Key Extraction From Sub-Addresses

179 Monero enables the creation of many sub-addresses from a single wallet, such that outputs to all addresses can be  
180 decrypted by the wallet’s main private view key  $k_s$ , but the subaddresses cannot be linked by a CA. Using Shor’s  
181 algorithm, keys can be extracted from sub-addresses as well.

182 The  $i$ th sub-address  $A_i$  contains the  $i$ th public spend key  $K_{s,i} := K_s \cdot g^{H(k_v,i)}$  and  $i$ th public view key  $K_{v,i} := K_{s,i}^{k_v}$ .  
183 Thus if an adversary learns of a subaddress, a QA can apply Shor’s algorithm to  $K_{v,i}$  to compute the discrete logarithm  
184 with respect to  $K_{s,i}$ , obtaining  $k_v$ . Using  $k_v$  the QA can classically compute  $H(k_v,i)$  for each allowable  $i$  in the core  
185 implementation. The adversary can then brute force compute  $K_{s,i} \cdot g^{-H(k_v,i)}$  through some small set of possible  
186 indexes  $i$  and obtain possible group elements, one of which is certain to be the  $K_{s,i}$  (in the core implementation).

187 From this, a CA can compute  $k_v$  (if  $k_v$  is deterministically chosen using  $k_s$  as in the core implementation) or a  
188 QA can compute  $k_v$  by inverting the map  $k_v \mapsto g^{k_v} = K_v$ .

### 189 2.1.3 Key Extraction From A Single One-Time Address

190 A single one-time address  $P$  is not sufficient to compute the private spend key  $k_s$ .

191 Key extraction from a single one-time key is not possible since one-time keys are perfectly hiding. Indeed, the  
192 one-time keys in Monero are of the form  $(R, P)$  for some transaction key  $R = g^r$  and for  $P = K_s \cdot g^{H((K_v)^r,i)}$ . The  
193 discrete logarithm, then, is  $k_s + H((K_v)^r,i)$ . Given any  $P$ , given any  $k_s$ , and given any  $k_v$ , there exists an  $r$  such  
194 that  $P = g^{k_s + H((K_v)^r,i)}$ .

195 In this way,  $P$  information-theoretically hides  $k_s$ . This means that, if only one output has ever been sent to  
196  $(K_s, K_v)$ , then at most 1 one-time key  $g^{k_s + H((K_v)^r,i)}$  appears on the blockchain. Even the QA, who can compute  
197 the discrete logarithm  $k_s + H((K_v)^r,i)$ , cannot determine  $k_s$  without additional information.

### 198 2.1.4 Key Extraction From A Pair of One-Time Addresses

199 In the previous section, we explained why a single one-time address is insufficient for key extraction. In this section,  
200 we describe how 2 one-time keys can be used to extract the spend key.

201 One of Monero’s classical security features is that addresses can be safely re-used, due to Monero’s one-time  
202 ‘stealth’ addresses, which prevent a CA from linking transactions to the same recipient or identifying the real  
203 address behind the stealth address.

204 However, in this sense, Monero is not secure against a QA: re-use of keys today allows a future hypothetical  
205 QA to extract the private spend key. If any address or subaddress (say with keypair  $(K_s, K_v)$ ) has received more  
206 than one transaction in the history of the blockchain, then the methods in this section can be applied by a QA with  
207 strictly public data to extract the keypair  $(k_v, k_s)$ .

208 Since creating a transaction is permissionless, this allows anybody with knowledge of your address to send multiple  
 209 outputs to your public keys, in the process making your private keys extractable to any QA at any time in the future.

210 Monero transactions are published with transaction keys  $R = g^r$  and one-time (so-called “stealth”) keys  $P =$   
 211  $K_s \cdot g^{H(K_v^r, i)}$ . Say that  $(R, P)$  and  $(R^*, P^*)$  are two pairs of keys for a Monero transaction made to the same address,  
 212 i.e.  $P = K_s \cdot g^{H((K_v)^r, i)}$  and  $P^* = K_s \cdot g^{H((K_v)^{r^*}, i^*)}$  for some  $i, i^*$ . A classical computer can compute  $P' = P \cdot (P^*)^{-1}$   
 213 easily. The QA can apply Shor’s algorithm to  $P, P^*$ , and  $P'$  obtaining the discrete logarithms  $p = k_s + H((K_v)^r, i)$ ,  
 214  $p^* = k_s + H((K_v)^{r^*}, i^*)$ , and  $p' = H((K_v)^r, i) - H((K_v)^{r^*}, i^*)$ , respectively. The QA can then classically compute  
 215  $p' + p_2 = p' - p_1 = k_s$ .

216 From this, a CA can compute  $k_v$  (if  $k_v$  is deterministically chosen using  $k_s$  as in the core implementation) or a  
 217 QA can compute  $k_v$  by inverting the map  $k_v \mapsto g^{k_v} = K_v$ .

## 218 2.2 Violate Signer Ambiguity

219 To prevent double-spends, Monero transactions require the publication of all images of the true signing keys used  
 220 in all ring signatures for the transaction under a one-way function. In this section, we show how the QA can use  
 221 a ring of public one-time keys and a key image known to be computed by one of the ring members to extract the  
 222 corresponding private one-time key.

223 For each transaction input, the signer includes a public key image  $J$  and a ring containing output one-time keys,  
 224 say  $\{P_1, \dots, P_n\}$ , where  $n$  is the ring size. At the time of this writing, it is necessary that all rings have exactly  
 225  $n = 11$  ring members to be considered valid. The message signer knows the private key  $p_\pi$  corresponding to some  
 226 public key  $P_\pi$  for one of the ring members, and learned the public keys for the other  $n - 1$  decoy keys by sampling  
 227 them from the blockchain.

228 Monero uses key images  $J := (H(P))^p$ . Under the discrete logarithm assumption, the CA cannot ascertain which  
 229 ring index  $\pi$  corresponds with the key  $P_\pi$  used by the signer to compute the key image  $J$ . The QA, on the other  
 230 hand, can compute  $H(P_i)$  for each  $i$  and compute the discrete logarithm  $\hat{p}_i$  of  $J$  with respect to each  $H(P_i)$ . For  
 231 some index  $\pi$ ,  $J = (H(g^{\hat{p}_\pi}))^{\hat{p}_\pi}$ . The QA concludes that the true signing key was  $P_\pi$  and has, in the course of drawing  
 232 this conclusion, learned  $p_\pi$ .

## 233 2.3 Violate Transaction Confidentiality

234 Since Pedersen commitments are perfectly hiding, and each transaction ostensibly uses fresh randomness, it should  
 235 seem that even the QA cannot look at an arbitrary Pedersen commitment to a transaction amount and compute the  
 236 transaction amount. However, the situation is not so nice.

### 237 2.3.1 Attack

238 The transaction amount is restricted to a small set  $[2^N]$  for some  $N \in \mathbb{N}$ . Recall that we commit to an amount  
 239  $b$  with a mask  $y$  by computing  $C(y, b) = g_1^y g_2^b$ , we see that there are at most  $2^N$  choices of  $b$ . This is selected to  
 240 be small enough so that users can brute-force search for the transaction amount when they see a new transaction  
 241 such that they can compute the blinder  $y$  (i.e. the transactions addressed to them). The QA can then apply Shor’s  
 242 algorithm  $2^N$  times to  $g_2^{-b^*} C(y, b)$  for each  $b^* \in [2^N]$  to obtain  $2^N$  possible choices of the blinder  $y$ . This allows the  
 243 QA to compute a Monero transaction amount with  $O(2^N)$  applications of Shor’s algorithm.

### 244 2.3.2 Mitigation

245 Since it is thought that the timeline for Shor’s algorithm to reliably compute the discrete logarithm for 32 byte keys  
 246 is even longer than the timeline for the arrival of the QA, one reasonable short-term mitigation for this violation  
 247 of privacy is to increase  $N$  substantially and to use some quantum-secure method of transmitting the transaction  
 248 amount, perhaps via secure sidechannel.

## 2.4 Violate Transaction Balancing

Monero ring confidential transactions in the style of [7] are inspired by Bitcoin-style confidential transactions from [9], swapping usual digital signatures for ring signatures. These transactions are proven to be balanced using range proofs of transaction amounts. In this section, we show how the QA can violate transaction balancing and we mention a mitigation.

### 2.4.1 Attack

Each Monero transaction is defined by some input anonymity sets (called rings) of possible input keys, some output keys, a non-negative plaintext fee, a range proof of the transaction amount including the fee, and a ring multisignature on the input rings. Monero amounts are encoded in perfectly hiding and computationally binding Pedersen commitments from [8], which are included as part of the input keys. A Monero transaction is considered valid if the anonymity sets are subsets of old output keys on the blockchain, the ring multisignature is valid (to authenticate whoever broadcast the transaction), and the range proof is valid (to ensure that transaction balances).

The Pedersen commitment scheme uses two basepoints,  $g_1$  and  $g_2$ , whose discrete logarithms are unknown with respect to each other. In the core Monero implementation,  $g_1 = g$  (the same as the public key basepoint) and  $g_2 = H(g)$  for some hash function  $H : \{0, 1\}^n \rightarrow G$ . To commit to an amount  $b$  with a mask  $y$ , we compute  $C(y, b) = g_1^y g_2^b$ .

Note that this map is many-to-one: for any  $(y, b)$ , there are many choices of  $(y', b') \neq (y, b)$  such that  $C(y', b') = C(y, b)$ . To open some  $C = C(y, b)$  to some  $(y', b') \neq (y, b)$ , a CA has few choices but to brute force search for any  $(y', b')$  that will do the trick. For a CA operating in PPT, this is considered intractable. In this sense, the Pedersen commitment scheme is computationally binding.

However, a QA can violate computational binding. The QA can apply Shor's algorithm to compute the discrete logarithm of  $g_2$  with respect to  $g_1$  or *vice versa*, say  $g_2 = g_1^\gamma = g^\gamma$ . To open  $C(y, b)$  to a different value  $b' \neq b$ , the QA can merely compute  $y' = y + \gamma(b - b')$  classically. Then  $C = C(y, b) = C(y', b')$  yet  $b \neq b'$ .

This resolves itself as the following attack. The QA receives a Monero output with amount  $b$ . The QA constructs (i) a range proof for use in a new transaction using amount  $b$  and the usual blinder as if they wish to spend all of  $b$ , but (ii) also uses the above approach to select a different blinder  $y'$  and a different amount  $b' < b$  for the recipient to open.

### 2.4.2 Mitigation

We identify two mitigations for the above attack.

First, if blinders are deterministically computed using random oracles from public data from the blockchain, then the QA is restricted from selecting  $y'$  freely. Depending on specific implementation recommendations, however, the computation for blinders in commitments may still be vulnerable to the QA, so we omit details.

Second, the idea of switch commitments are introduced in [11]. These commitments are homomorphic, so they can be used (similarly to Pedersen commitments) to the degree necessary for usage in a currency protocol. They already contain usual Pedersen commitments, and so they can be considered an extension of the commitment scheme Monero already uses.

These commitments have the following properties. An opener can convince a verifier that some  $C$  commits to a certain value with either a *partial opening* or a *full opening*. If the opener partially opens a commitment, then the scheme is computationally binding and perfectly hiding. If the opener fully opens a commitment, then the scheme is statistically binding and computationally hiding.

Switch commitments in Monero would require (i) a single additional group element in transactions, (ii) a slight modification to range proofs resulting in no significant change to transaction size or verification time, and (iii) the computation of an additional scalar by the recipient to open commitments.

292 Note that when the QA violates the binding property of commitments, they can violate the supply of Monero,  
293 which is intended to be a money with a known, fixed supply secure against malicious tampering. On the other  
294 hand, when the QA violates the hiding property, they can violate the amount privacy afforded by the confidential  
295 transactions of Monero.

296 A protocol that does not balance transactions cannot function as a currency, whether transactions are confidential  
297 or not. A currency without confidential transactions can still function as a currency, however. In this way, we could  
298 regard the monetary supply as hierarchically more important to secure than user privacy.

## 299 2.5 Violate Unlinkability

300 The integration of one-time addresses into Monero’s functionality has provided transactions with an extra layer of  
301 protection. In their essence, one-time stealth addresses work by masking the public keys used in each transaction.  
302 These one-time stealth addresses can be formed either from subaddresses or the primary address. Using a combination  
303 of Shor’s and Grover’s algorithm, a quantum computer could reveal the genuine public keys behind each of the stealth  
304 addresses being used in a transaction.

305 Since this mechanism relies on a QA using Grover’s algorithm to find the pre-image of a hash digest used to  
306 generate the stealth address from spend and view keys, it is somewhat more robust against a QA compared to the  
307 other vulnerabilities mentioned, so the trick is to narrow down possible pre-image values to decrease the state space  
308 needed to run iterations of Grover’s algorithm. For a 256 bit hash digest, even for a quantum computer to find the  
309 pre-image using a pure brute force attack this would require approximately  $2^{128}$  iterations of Grover’s algorithm to  
310 output a corresponding amplitude, which at standard clock speeds would require a runtime much greater than the  
311 lifetime of the universe.

312 If Grover’s algorithm was instead used to go through all known public view/spend key pairs until it found a  
313 match that output the stealth address with a known  $r$  (from Shor’s) or perhaps used to go through a smaller state  
314 space somehow (methods shall be described below) then this functionality is vulnerable to a QA.

315 If all transactions made by a public key are generated off a new subaddress that is different each time, this could  
316 provide an extra layer of security that might make it much more difficult

317 If Alice wants to send a transaction to Bob, a one time stealth address is formed as such: First, generating a  
318 random integer  $r$  to be used only once in this transaction,  $K_0 = \mathcal{H}_n(rK_v)G + K_s$  Where  $K_0$  is the one-time stealth  
319 address Alice submits this and  $g^r$  (the transaction public key) to the network. Once submitted to the network, Bob  
320 can see it is meant for him by multiplying the transaction public key using his private view key  $k_v$ , since  $rk_vG = rK_v$ .  
321 From this,  $K'_s = K_0 - \mathcal{H}_n(rK_v)G = K_s$ . This is considered secure classically, however a quantum adversary could  
322 break this feature as such: Find  $r$ , by factoring  $g^r$  using Shor’s. Using Grover’s, find a  $K$  spend and view public key  
323 pair amongst existing public keys such that  $K_0 = \mathcal{H}_n(rK_v)G + K_s$  for a given  $r$ . From this, extracting private keys  
324 is simple, using Shor’s to invert  $k \mapsto g^k = K$  as described in section 2.1.1.

325 If a subaddress was used to generate the stealth address instead of a public key address, then this would add an  
326 extra layer of security. However, since these are also secured by DLP, it would be possible to find the associated  
327 public key address used to generate the relevant subaddress attached to the relevant stealth address from this in  
328 much the same manner as shown.

## 329 2.6 Decrypt payment identifiers

330 Monero transactions optionally come with payment identification that consist of the XOR between a bitstring message  
331 and a mask. The mask is generated from the hash of  $K_v^r$ , where the transaction key is  $R = g^r$ . Thus, in any of the  
332 previous sections where the QA can compute  $k_v$  using Shor’s algorithm, the QA can apply Shor’s algorithm a second  
333 time to  $R$  to compute  $r$  and compute the mask  $X$  directly.



## 334 2.7 Immutability

### 335 2.7.1 Block immutability

336 M question - can Grover's algorithm be used to forge a block with target hash and desired arbitrary payload?

337 Can you circumvent the pq-PoW hardness of RandomX by brute forcing other parts of the block (for example  
338 tx\_extra in the miner coinbase)

339 I have a bit more clarity now about how this might work.

340 (here, we answer the question: if in the future RandomX is the pq-bottleneck and you can Grover's effectively,  
341 what are your options?)

342 The attacker actually do a proof of work! Investing the full amount of time/energy with with a classical computer

343 So if you a bunch of block information  $B$  and you boil it down to a digest  $b$  to hash with RandomX, so we try a  
344 bunch of nonces:

345  $RandomX(b, n_1) \rightarrow$  fail to pass difficulty

346  $RandomX(b, n_2) \rightarrow$  fail to pass difficulty

347  $RandomX(b, n_3) \rightarrow$  fail to pass difficulty

348  $RandomX(b, n_4) \rightarrow$  fail to pass difficulty

349  $RandomX(b, n_5) \rightarrow$  fail passes difficulty ... winning block

350 Now, I don't broadcast anything. I just stick  $(b, n_5)$  in my back pocket as winning inputs to RandomX. Then,  
351 whenever I feel like it, I use Grover's algorithm to find a  $B$  preimage for the  $b$  that I previously solved.

352 You don't have to use QA tricks on RandomX if you can pull off QA tricks on its inputs.

353 Actually, heh, I don't think you have to do any proof of work at all! You could just pick random previously solved  
354 blocks, take their  $(n, b)$  pair, and make your own  $B \rightarrow b$  mapping with Grovers. (of course that would be sloppy,  
355 and Noncesense would eventually notice the RandomX nonce collision and start poking around.)

356 Does this make any sense?

### 357 2.7.2 Transaction immutability

358 M - will finish this up later notes - Txns have no PoW

359 Grover's to find primage of alternate txn that produces the same txid

360 Fluffy blocks are a risk if QAs are forging transactions. More broadly, the fact that the block PoW input only  
361 ingests transaction IDs instead of transaction data makes this an attack surface

## 362 **3 Alternatives to Elliptic Curve Based Cryptography**

### 363 **3.1 Lattice-based**

#### 364 **3.1.1 From a geometric point of view**

#### 365 **3.1.2 Differences between lattice-based cryptography, RSA, and ECC**

#### 366 **3.1.3 Migration**

### 367 **3.2 Multivariate-based**

### 368 **3.3 Hash-based**

### 369 **3.4 Supersingular Elliptic Curve Isogeny-Based Cryptography**

## 370 **4 Alternative protocols**

### 371 **4.1 ZK-STARKS**

#### 372 **4.1.1 Possible usage**

### 373 **4.2 MatRiCT**

### 374 **4.3 L2RS**

### 375 **4.4 RingRainbow**

## 376 **5 Key Size/Verification Time Table**

## 377 **6 Conclusion**

## 378 **A Technical Implementation of Quantum Algorithms**

379 The introduction to this technical note outlines some of the capabilities a theoretical quantum adversary could  
380 leverage in order to introduce vulnerabilities into the current Monero security infrastructure. This appendix serves  
381 as a more in depth technical look into the algorithms that can be run efficiently on quantum hardware in order  
382 to provide these capabilities. The three algorithms provided in this appendix include Shor's, Grover's and Simon's  
383 algorithm. While these three do provide some of the largest security risks to Monero as outlined, the capabilities of  
384 a quantum computer are not limited to them. The main security concern not included in this group might be from  
385 algorithms that can be used in QDC (Quantum Differential Cryptanalysis) such as the Bernstein-Vazirani algorithm  
386 and others. The main threat factor from QDC would be in reversing or revealing hidden information about the  
387 pre-image of a hash digest.

### 388 **1. A.1 Shor's Algorithm Technical Implementation**

389 Shor's algorithm violates the RSA hardness assumption [3] and the elliptic curve discrete logarithm hardness  
390 assumption [10], both with polynomial time complexity.

391 Given that some of the largest security assumptions of which Monero is based off of can be solved efficiently  
392 using a specialized elliptic curve implementation of Shor's algorithm, the specific implementation will be shown  
393 as such:

394 Elliptic curves over finite fields form abelian groups. For discrete logarithms over elliptic curves, We consider  
 395 an elliptic curve,  $E$ , over  $GF(p)$ , where  $p$  is a large prime. The base of the logarithm is a point  $P$  on the  
 396 elliptic curve  $E$ , whose order is another (large) prime  $q$ , such that  $qP = 0$ . We want to compute the discrete  
 397 logarithm,  $d$ , that lies on another point  $Q$  on the elliptic curve such that  $Q = dP$ .

398 We apply the following transformation to a set of qubits, such that the quantum state can be described as:

$$|\psi\rangle = \frac{1}{2^n} \sum_{x=0}^{2^n-1} \sum_{y=0}^{2^n-1} |x, y, xP + yQ\rangle \quad (1)$$

399 Thus we need a method of computing (large) integer multiples of group elements. This can be done by the  
 400 standard “double and add technique”. This is the same technique used for the modular exponentiation in the  
 401 factoring algorithm, although there the group is written multiplicatively so it’s called the square and multiply  
 402 technique. To compute  $xP + yQ$ , first we repeatedly double the group elements  $P$  and  $Q$ , thus getting the  
 403 multiples  $P_i = 2^i P$  and  $Q_i = 2^i Q$ . We then add together the  $P_i$  and  $Q_i$  for which the corresponding bits of  $x$   
 404 and  $y$  are 1,  $xP + yQ$  can then be written as

$$xP + yQ = \sum_i x_i P_i + \sum_i y_i Q_i \quad (2)$$

405 The multiples as such of  $P_i$  and  $Q_i$  can be computed classically beforehand, creating conditions that can be  
 406 used describe the state shown in equation (2) for  $|x, y\rangle$  using just a single qubit, which drastically cuts down  
 407 the qubit requirements. This is done by using the semiclassical quantum Fourier transform [2] and is analogous  
 408 to the factoring algorithm used in breaking RSA. Thus, we can represent  $|x, y, xP + yQ\rangle = |x, y, O\rangle$  with  $O$   
 409 simply being the "accumulator" register, with  $|x, y\rangle$  being replaced by the single qubit as mentioned before.

410 We are left being required to carry out a number of steps whereby we add a fixed (classically known) point  
 411  $P_i$  (or  $Q_i$ ) to a superposition of points. We are working in the cyclic group generated by  $P$ , thus the effect of  
 412 a fixed addition is to “shift” the discrete logarithm of each element in the superposition by the same amount.  
 413 Thus we need unitary transformations  $U_{P_i}$  and  $U_{Q_i}$  which acts on any basis state  $|S_i\rangle$  representing a point on  
 414 the elliptic curve, as  $U_{P_i} : |S\rangle \rightarrow |S + P_i\rangle$  and  $U_{Q_i} : |S\rangle \rightarrow |S + Q_i\rangle$ .

415 Applying these steps  $n$  times to  $P$  and  $n$  times to  $Q$  where  $n$  is approximately  $n = \log_2 q$ . This sequence of steps  
 416 decomposes the discrete logarithm quantum algorithm into a sequence of group shifts by constant classically  
 417 known elements. After these are done, the resultant measured output will correspond to a distribution with  
 418 peaks around points equal to  $Nk/q$  and  $Ndk/q$  where  $N = 2^n$ . To obtain the relevant values,  $k$  and  $dk$  which  
 419 are ultimately what will be used to solve the DLP, we multiply the observed output values by  $q/N$ .

420 This version of Shor’s algorithm for solving the DLP for elliptic curves for 256 bit security can be done efficiently  
 421 at standard clock speeds in under a minute, which for a classical computer using the most efficient currently  
 422 known algorithm would take over ten-thousand years to find the same result. As such, it should be obvious  
 423 why this is a considerable advantage a QA (Quantum Adversary) would have over a CA (Classical Adversary).

## 424 2. A.2 Grover’s Algorithm Technical Implementation

425 Given a set of data, a targeted implementation of Grover’s Algorithm can procure from this data a specific  
 426 labeled instance  $x_0$  which can be for example, the confidential pre-image of a hash digest. These are the steps  
 427 one would take to do so:

428 (a) Initialize the qubits

$$|\psi\rangle = |0\rangle^{\otimes n}$$

429 (b) Put the qubits in an equal state of superposition.

$$|s\rangle = H^{\otimes n}|0\rangle^n$$

430 (c) Apply an oracle reflection  $|U_f\rangle$  to the marked instance  $x_0$  of the qubits.

431 (d) Apply an additional reflection,

$$U_s = 2|s\rangle\langle s| - I$$

432 Such that this maps the state to  $U_s U_f |s\rangle$  Repeat steps 3-4 approximately  $\sqrt{N} = t$  times, where  $N$  is the  
433 number of entries in the data set.

434 (e) Once the state of the system can be described as  $|\psi_t\rangle = (U_s U_f)^t |s\rangle$ , you measure the qubits, the corre-  
435 sponding amplitude will correspond to the equivalent classical bits of the target entry.

### 436 3. A.3 Simon's Algorithm Technical Implementation

437 Simon's problem is, given a function  $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$  that is known to be invariant under some  $n$ -bit XOR  
438 mask  $a$ , determine  $a$ . In other words, given  $f(x) = f(y)$  if and only if  $x \oplus y \in \{0, a\}$ , compute  $a$ .

439 This can be used to set up a system of linear equations that can be used to find the outputs and find the XOR  
440 mask of certain functions, which has numerous applications in cryptography.

441 1. first you initialize two  $n$ -qubit registers to 0

$$|\psi_1\rangle = |0\rangle^{\otimes n} |0\rangle^{\otimes n} \quad (3)$$

442 2. Apply a Hadamard transform to the first register

$$|\psi_2\rangle = \frac{1}{\sqrt{2^n}} \sum_{x \in \{0,1\}^n} |x\rangle |0\rangle^{\otimes n} \quad (4)$$

443 3. Apply a query function  $Q_f$

$$|\psi_3\rangle = \frac{1}{\sqrt{2^n}} \sum_{x \in \{0,1\}^n} |x\rangle |f(x)\rangle \quad (5)$$

444 4. Measure second register, causing the first register to become

$$|\psi_4\rangle = \frac{1}{\sqrt{2}} (|x\rangle + |y\rangle) \quad (6)$$

445 5. Apply Hadamard transform to the first register

$$|\psi_5\rangle = \frac{1}{\sqrt{2^{n+1}}} \sum_{z \in \{0,1\}^n} [(-1)^{x \cdot z} + (-1)^{y \cdot z}] |z\rangle \quad (7)$$

446 6. Measure the first register, which gives an output if

$$(-1)^{x \cdot z} = (-1)^{y \cdot z} \quad (8)$$

447 The output will correspond to a string  $z$ , this string will correspond to  $b \cdot z = 0 \pmod{2}$  which from Gaussian  
448 elimination can be used to find the XOR mask to the function  $f(x)$ . This can be run exponentially faster than  
449 any equivalent classical algorithm.

## 450 References

- 451 [1] Charles H Bennett, Ethan Bernstein, Gilles Brassard, and Umesh Vazirani. Strengths and weaknesses of quantum  
452 computing. *SIAM journal on Computing*, 26(5):1510–1523, 1997.
- 453 [2] Richard Cleve, Artur Ekert, Chiara Macchiavello, and Michele Mosca. Quantum algorithms revisited. *Proceed-*  
454 *ings of the Royal Society of London. Series A: Mathematical, Physical and Engineering Sciences*, 454(1969):339–  
455 354, 1998.
- 456 [3] Edward Gerjuoy. Shor’s factoring algorithm and modern cryptography. an illustration of the capabilities inherent  
457 in quantum computers. *American journal of physics*, 73(6):521–540, 2005.
- 458 [4] Lov K Grover. A fast quantum mechanical algorithm for database search. In *Proceedings of the twenty-eighth*  
459 *annual ACM symposium on Theory of computing*, pages 212–219, 1996.
- 460 [5] Pascal Koiran, Vincent Nesme, and Natacha Portier. A quantum lower bound for the query complexity of  
461 simon’s problem. In *International Colloquium on Automata, Languages, and Programming*, pages 1287–1298.  
462 Springer, 2005.
- 463 [6] Michele Mosca. Quantum algorithms. *arXiv preprint arXiv:0808.0369*, 2008.
- 464 [7] Shen Noether, Adam Mackenzie, et al. Ring confidential transactions. *Ledger*, 1:1–18, 2016.
- 465 [8] Torben Pryds Pedersen. Non-interactive and information-theoretic secure verifiable secret sharing. In *Annual*  
466 *international cryptology conference*, pages 129–140. Springer, 1991.
- 467 [9] Andrew Poelstra, Adam Back, Mark Friedenbach, Gregory Maxwell, and Pieter Wuille. Confidential assets. In  
468 *International Conference on Financial Cryptography and Data Security*, pages 43–63. Springer, 2018.
- 469 [10] John Proos and Christof Zalka. Shor’s discrete logarithm quantum algorithm for elliptic curves. *arXiv preprint*  
470 *quant-ph/0301141*, 2003.
- 471 [11] Tim Ruffing and Giulio Malavolta. Switch commitments: A safety switch for confidential transactions. In  
472 *International Conference on Financial Cryptography and Data Security*, pages 170–181. Springer, 2017.
- 473 [12] Thomas Santoli and Christian Schaffner. Using simon’s algorithm to attack symmetric-key cryptographic prim-  
474 itives. *arXiv preprint arXiv:1603.07856*, 2016.
- 475 [13] Daniel R Simon. On the power of quantum computation. *SIAM journal on computing*, 26(5):1474–1483, 1997.
- 476 [14] Huiqin Xie and Li Yang. Using bernstein–vazirani algorithm to attack block ciphers. *Designs, Codes and*  
477 *Cryptography*, 87(5):1161–1182, 2019.